

RANJIT JHALA, ERIC SEIDEL, NIKI VAZOU,
EDITED BY CATARINA GAMBOA

PROGRAMMING WITH REFINEMENT TYPES

AN INTRODUCTION TO LIQUIDHASKELL

Version 13, July 20th, 2020, edit on Mar, 2024.

Copyright © 2024 Ranjit Jhala

[HTTPS://UCSD-PROGSYS.GITHUB.IO/LIQUIDHASKELL-BLOG/](https://ucsd-progsys.github.io/liquidhaskell-blog/)

Licensed under the Apache License, Version 2.0 (the “License”); you may not use this file except in compliance with the License. You may obtain a copy of the License at <http://www.apache.org/licenses/LICENSE-2.0>. Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an “AS IS” BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

Contents

1	<i>Introduction</i>	7
	<i>Well-Typed Programs Do Go Wrong</i>	7
	<i>Refinement Types</i>	8
2	<i>Refinement Types</i>	11
	<i>Defining Types</i>	11
	<i>Errors</i>	12
	<i>Subtyping</i>	13
	<i>Writing Specifications</i>	14
	<i>Refining Function Types: Post-conditions</i>	14
	<i>Dependent Refinements</i>	15
3	<i>Refined Datatypes</i>	17
	<i>Sparse Vectors</i>	18
	<i>Queues</i>	21
	<i>Sized Lists</i>	23
	<i>Queue Type</i>	25
	<i>Queue Operations</i>	26
	<i>Recap of everything!</i>	29

4	<i>Cheat Sheet</i>	31
	<i>Specifications</i>	31
	<i>Alias</i>	31
	<i>Liquid types in Datatypes</i>	31
	<i>Measures</i>	32
5	<i>Refined Datatypes</i>	33
	<i>Sparse Vectors</i>	34

List of Exercises

3.1	Exercise (Destructing Lists)	25
3.2	Exercise (Queue Sizes)	26
3.3	Exercise (Insert)	28
3.4	Exercise (Rotate)	29

1

Introduction

Welcome to the LiquidHaskell Short Tutorial, where you will learn the basic workings of LiquidHaskell and complete some exercises. The full version of the tutorial can be found in the [project's website](#).

One of the great things about Haskell is its brainy type system that allows one to enforce a variety of invariants at compile time, thereby nipping in the bud a large swathe of run-time errors.

Well-Typed Programs Do Go Wrong

Alas, well-typed programs *do* go quite wrong, in a variety of ways.

DIVISION BY ZERO This innocuous function computes the average of a list of integers:

```
average    :: [Int] -> Int
average xs = sum xs `div` length xs
```

We get the desired result on a non-empty list of numbers:

```
ghci> average [10, 20, 30, 40]
25
```

However, this program crashes with certain arguments. From the following options, what argument would make average crash?

[1] [] [1,1,1,1,1,1,1,1,1,1] Submit

Answer

If we call it with an empty list, we get a rather unpleasant crash:
 *** Exception: divide by zero. We could write average more *defensively*, returning a Maybe or Either value. However, this merely kicks the can down the road. Ultimately, we will want to extract the Int from the Maybe and if the inputs were invalid to start with, then at that point we'd be stuck.

HEART BLEEDS

For certain kinds of programs, there is a fate worse than death. `text` is a high-performance string processing library for Haskell, that is used, for example, to build web services.

```
ghci> :m +Data.Text Data.Text.Unsafe
ghci> let t = pack "Voltage"
ghci> takeWord16 5 t
"Volta"
```

A cunning adversary can use invalid, or rather, *well-crafted*, inputs that go well outside the size of the given text to read extra bytes and thus *extract secrets* without anyone being any the wiser.

```
ghci> takeWord16 20 t
"Voltage\1912\3148\SOH\NUL\15928\2486\SOH\NUL"
```

The above call returns the bytes residing in memory *immediately after* the string `Voltage`. These bytes could be junk, or could be either the name of your favorite TV show, or, more worryingly, your bank account password.

Refinement Types

Refinement types allow us to enrich Haskell's type system with *predicates* that precisely describe the sets of *valid* inputs and outputs of functions, values held inside containers, and so on. These predicates are drawn from special *logics* for which there are fast *decision procedures* called SMT solvers.

BY COMBINING TYPES WITH PREDICATES you can specify *contracts* which describe valid inputs and outputs of functions. The refinement type system *guarantees at compile-time* that functions adhere to their contracts. That is, you can rest assured that the above calamities *cannot occur at run-time*.

LIQUIDHASKELL is a Refinement Type Checker for Haskell, and in this tutorial we'll describe how you can use it to make programs better and programming even more fun.

As a glimpse of what LiquidHaskell can do, run the average example below by pushing the green triangle on the top, and try to read the error message. Since `div` cannot take a zero value as the second argument, and LiquidHaskell sees that it is a possibility in this function, an error will be raised.

```
average'    :: [Int] -> Int
average' xs = sum xs `div` length xs
```

In this tutorial you will learn how to add and reason about refinement types in Haskell, and how it can increase the reliability of Haskell problems.

Next

2

Refinement Types

WHAT IS A REFINEMENT TYPE? In a nutshell,

$$\text{Refinement Types} = \text{Types} + \text{Predicates}$$

That is, refinement types allow us to decorate types with *logical predicates*, which you can think of as *boolean-valued* Haskell expressions, that constrain the set of values described by the type. This lets us specify sophisticated invariants of the underlying values.

Defining Types

Let us define some refinement types:

```
{-@ type Zero    = {v:Int | v == 0} @-}  
{-@ type NonZero = {v:Int | v /= 0} @-}
```

THE VALUE VARIABLE v denotes the set of valid inhabitants of each refinement type. Hence, `Zero` describes the *set of* `Int` values that are equal to `0`, that is, the singleton set containing just `0`, and `NonZero` describes the set of `Int` values that are *not* equal to `0`, that is, the set `{1, -1, 2, -2, ...}` and so on.

To indicate that these specifications are for `LiquidHaskell` we write them like `{-@ spec @-}`.

Now, TO USE these types we can write:

```

{-@ zero :: Zero @-}
zero = 0 :: Int

{-@ one, two, three :: NonZero @-}
one   = 1 :: Int
two   = 2 :: Int
three = 3 :: Int

```

Errors

If we try to say nonsensical things like:

```

nonsense :: Int
nonsense = one'
  where
    {-@ one' :: Zero @-}
    one' = 1

```

LiquidHaskell will complain with an error message:

```
../liquidhaskell-tutorial/src/03-basic.lhs:72:3-6: Error: Liquid Type Mismatch
```

```

72 |   one' = 1 :: Int
    |     ^^^^

Inferred type
  VV : {VV : Int | VV == (1 : int)}

not a subtype of Required type
  VV : {VV : Int | VV == 0}

```

The message says that the expression `1 :: Int` has the type

```
{v:Int | v == 1}
```

which is *not* (a subtype of) the *required* type

```
{v:Int | v == 0}
```

as 1 is not equal to 0.

Subtyping

What is this business of *subtyping*? Suppose we have some more refinements of `Int`

```
{-@ type Nat      = {v:Int | 0 <= v}      @-}

{-@ type Positive = {v:Int | 0 < v}        @-}

{-@ type Even     = {v:Int | v mod 2 == 0 } @-}

{-@ type TensToHundred = {v:Int | v mod 10 == 0 && v <= 100} @-}
```

SUBTYPING AND IMPLICATION

Zero is the most precise type for `0 :: Int`, as it is a *subtype* of `Nat`, `Even`. However, it is not a subtype of `Positive`. The alias `TensToHundred` represents the multiples of 10 smaller than 100, meaning that `Even` is a *subtype* of it but all the other ones are not.

Exercise: Now let us try a new predicate.

Write a type `Percentage` for the numbers that represent a percentage (between 0 and 100).

Then, remove the comment the liquid type signatures and run the code. The first example should be correct and the second should not.

```
-- write the alias here

-- {-@ percentT  :: Percentage @-}
percentT      = 10 :: Int
-- {-@ percentF  :: Percentage @-}
percentF      = 10 + 99 :: Int
```

Answer

```
{-@ type Percentage = {v:Int | 0 <= v && v <= 100} @-}
```

IN SUMMARY the key points about refinement types are:

1. A refinement type is just a type *decorated* with logical predicates.
2. A term can have *different* refinements for different properties.
3. When we *erase* the predicates we get the standard Haskell types.

Writing Specifications

We can also add specifications as pre- and post-conditions of functions.

Remember the divide function from before? We can add the case of dividing by zero with this die "message" to indicate that this case should be handled before running the code.

```
divide'      :: Int -> Int -> Int
divide' n 0 = die "divide by zero"
divide' n d = n `div` d
```

So, now we can specify that the first case will never with a *pre-condition* that says that the second argument is non-zero:

```
{-@ divide :: Int -> NonZero -> Int @-}
divide _ 0 = die "divide by zero"
divide n d = n `div` d
```

You can run the both pieces of code and check that the first one throws an error while the second one does not since it can infer that the first case will not be called.

ESTABLISHING PRE-CONDITIONS

The above signature forces us to ensure that that when we *use* divide, we only supply provably NonZero arguments.

Exercise:

Select which of the following functions that call divide would raise an error:

```
foo x y = divide (x + y) 2
foo' x y z = divide (divide (x + y) 3) 10
foo'' x y z = divide (x + y) z
```

Submit

Answer

foo'' is the invocation that could trigger a crash since we have no guarantees that z is a NonZero value.

Refining Function Types: Post-conditions

Next, let's see how we can use refinements to describe the *outputs* of a function. Consider the following simple *absolute value* function

```
abs          :: Int -> Int
abs n
  | 0 < n    = n
  | otherwise = 0 - n
```

We can use a refinement on the output type to specify that the function returns non-negative values

```
{-@ abs :: Int -> Nat @-}
```

LiquidHaskell *verifies* that `abs` indeed enjoys the above type by deducing that `n` is trivially non-negative when $0 < n$ and that in the otherwise case, the value $0 - n$ is indeed non-negative.

Dependent Refinements

The predicates in pre- and post- conditions can also refer to previous arguments of the function.

For example, including that the output is greater than the input.

```
{-@ plus1 :: a:Int -> {b:Int | b > a}@-}
plus1 :: Int -> Int
plus1 a = a + 1
```

And the same could be done between input values.

Exercise:

Let's put everything together now.

Write a specification for the method `calcPer` that:

- 1) first receives a positive int;
- 2) then an int with a value between zero and the first int;
- 3) returns a percentage;

Use the aliases created in the exercises you have completed before.

```
calcPer      :: Int -> Int -> Int
calcPer a b  = (b * 100) `div` a

cpc = calcPer 10 5 :: Int -- should be correct
cpi = calcPer 10 11 :: Int -- should be incorrect
```

Answer

```
{-@ calcPer :: a:Positive -> {b:Int | 0 <= b && b <= a} ->  
c:Percentage @-}
```

You finished the first part of the Tutorial! Tell the interviewers you got to the end of the page, and answer some questions from our team before moving to the next section.

Next

3

Refined Datatypes

So far, we have seen how to refine the types of *functions*, to specify, for example, pre-conditions on the inputs, or post-conditions on the outputs. In this section we will see how to apply these in datatypes. First, by defining properties of data values, and then by defining *datatypes* that satisfy certain invariants. In the latter, it is handy to be able to directly refine the data definition, making it impossible to create illegal inhabitants.

MEASURES are used to define *properties* of Haskell data values that are useful for specification and verification.

A MEASURE is a *total* Haskell function, 1. With a *single* equation per data constructor, and 2. Guaranteed to *terminate*, typically via structural recursion.

We can tell LiquidHaskell to *lift* a function meeting the above requirements into the refinement logic by declaring:

```
{-@ measure nameOfMeasure @-}
```

For example, for a list we can define a way to *measure* its size with the following function.

```
{-@ measure size @-}  
{-@ size :: [a] -> Nat @-}  
size :: [a] -> Int  
size []      = 0  
size (_:rs) = 1 + size rs
```

Then, we can use this measure to define aliases.

Let's create another measure named `notEmpty` that takes a list as input and returns a `Bool` with the information if it is empty or not.

```
-- write notEmpty measure
```

Answer

```
{-@ measure notEmpty @-} notEmpty      :: [a] -> Bool
notEmpty []      = False notEmpty (_,_) = True
```

We can now define a couple of useful aliases for describing lists of a given dimension.

For example, we can define that a list has exactly N elements.

```
{-@ type ListN a N = {v:[a] | size v == N} @-}
```

Note that when defining refinement type aliases, we use uppercase variables like N to distinguish *value* parameters from the lowercase *type* parameters like a .

Now, try to create an alias `NEList` for an empty list, using the measure `notEmpty` created before. When removed from comment, the first example should raise an error while the second should not.

```
-- write the alias here

-- Remove the comments below to test the alias
-- {-@ ne1 :: NEList Int@-}
-- ne1 = [] :: [Int]      -- should fail
-- {-@ ne2 :: NEList Int@-} -- should be correct
-- ne2 = [1,2,3,4] :: [Int]
```

Answer

```
{-@ type NEList a = {v:[a] | notEmpty v} @-}
```

Sparse Vectors

As our first example of a refined datatype, let's see Sparse Vectors. While the standard `Vector` is great for dense arrays, often we have to manipulate sparse vectors where most elements are just 0. We might represent such vectors as a list of index-value tuples `[(Int, a)]`.

Let's create a new datatype to represent such vectors:

```
data Sparse a = SP { spDim    :: Int
                    , spElems :: [(Int, a)] }
```

Thus, a sparse vector is a pair of a dimension and a list of index-value tuples. Implicitly, all indices *other* than those in the list have the value 0 or the equivalent value type `a`.

LEGAL

Sparse vectors satisfy two crucial properties. 1. the dimension stored in `spDim` is non-negative;

2. every index in `spElems` must be valid, i.e. between 0 and the dimension.

Unfortunately, Haskell's type system does not make it easy to ensure that *illegal vectors are not representable*.

DATA INVARIANTS LiquidHaskell lets us enforce these invariants with a refined data definition:

```
{-@ data Sparse a = SP { spDim    :: Nat
                        , spElems :: [(Btwn 0 spDim, a)] } @-}
```

Where, as before, we use the aliases:

```
{-@ type Nat      = {v:Int | 0 <= v}      @-}
{-@ type Btwn Lo Hi = {v:Int | Lo <= v && v < Hi} @-}
```

REFINED DATA CONSTRUCTORS The refined data definition is internally converted into refined types for the data constructor `SP`. So, by using refined input types for `SP` we have automatically converted it into a *smart* constructor that ensures that *every* instance of a `Sparse` is legal. Consequently, LiquidHaskell verifies:

```
okSP :: Sparse String
okSP = SP 5 [ (0, "cat")
              , (3, "dog") ]
```

but rejects, due to the invalid index:

```
badSP :: Sparse String
badSP = SP 5 [ (0, "cat")
               , (6, "dog") ]
```

Write another example of a Sparse data type that is invalid. Remove the comment from the type signature below and complete the implementation with the example.

```
-- badSP' :: Sparse String
```

Answer

e.g., `badSP' = SP (-1) [(0, "cat")]`

FIELD MEASURES It is convenient to write an alias for sparse vectors of a given size N . So that we can easily say in a refinement that we have a sparse vector of a certain size.

For this we can use *measures*.

MEASURES WITH SPARSE VECTORS

Similarly, the sparse vector also has a *measure* for its dimension, and it is already defined by `spDim`, so we can use it to create the new alias of sparse vectors of size N .

Think Aloud:

For the following exercise, we will use a technique called Think Aloud, where you should try to say everything that comes to your mind while you engage with the exercise.

In specific, aim:

- a) to speak all thoughts, even if they are unrelated to the task;
- b) to refrain from explaining the thoughts;
- c) to not try to plan out what to say;
- d) to imagine that you are alone and speaking to yourself; and
- e) to speak continuously.

For the following exercise, read the question aloud and remember to voice your thoughts while solving the exercise.

Following what we did with the lists, write the alias `SparseN` for sparse vector of length N , using `spDim` instead of `size`.

```
-- write the alias here
```

Answer

e.g., `{-@ type SparseN a N = {v: Sparse a | spDim v == N} @-}`

YOU FINISHED THE TUTORIAL!

You finished the Tutorial! Tell the interviewers you got to the end of the page, and answer some questions from our team before moving to the next section.

NEXT EXERCISE!

Now that you have learned the main blocks of LiquidHaskell, let's complete an exercise using all the concepts.

You can open a Cheat Sheet with examples of the main concepts on another tab on the side.

Next Case Study: Okasaki's Lazy Queues =====

Lets test what we learned so far in a case study that is simple enough to explain without pages of code, yet complex enough to show off whats cool about dependency: Chris Okasaki's beautiful [Lazy Queues](#). This structure leans heavily on an invariant to provide fast *insertion* and *deletion*. Let's see how to enforce that invariant with LiquidHaskell.

Queues

A [queue](#) is a structure into which we can insert and remove data such that the order in which the data is removed is the same as the order in which it was inserted.




Figure 3.1: A Queue is a structure into which we can insert and remove elements. The order in which the elements are removed is the same as the order in which they were inserted.

TO EFFICIENTLY IMPLEMENT a queue we need to have rapid access to both the front as well as the back because we remove elements from former and insert elements into the latter. This is quite straightforward with explicit pointers and mutation – one uses an old school linked list and maintains pointers to the head and the tail. But can we implement the structure efficiently without having stoop so low?

CHRIS OKASAKI came up with a very cunning way to implement queues using a *pair* of lists – let's call them front and back which represent the corresponding parts of the Queue.

- To insert elements, we just *cons* them onto the back list,
- To remove elements, we just *un-cons* them from the front list.




Figure 3.2: We can implement a Queue with a pair of lists; respectively representing the front and back.

THE CATCH is that we need to shunt elements from the back to the front every so often, e.g. we can transfer the elements from the back to the front, when:

1. a remove call is triggered, and
2. the front list is empty.

OKASAKI'S FIRST INSIGHT was to note that every element is only moved *once* from the back to the front; hence, the time for insert and remove could be $O(1)$ when *amortized* over all the operations. This is perfect, *except* that some set of unlucky remove calls (which




Figure 3.3: Transferring Elements from back to front.

occur when the front is empty) are stuck paying the bill. They have a rather high latency up to $O(n)$ where n is the total number of operations.

OKASAKI'S SECOND INSIGHT saves the day: he observed that all we need to do is to enforce a simple *balance invariant*:

$$\text{Size of front} \geq \text{Size of back}$$

THIS IS A GOOD MOMENT TO SEE IF YOU UNDERSTOOD THE IDEA OF THESE QUEUES, WE WILL START IMPLEMENT THEM NOW.

LET'S IMPLEMENT QUEUES and ensure the crucial invariant(s) with LiquidHaskell. What we need are the following ingredients:

1. A type for Lists, and a way to track their size,
2. A type for Queues which encodes the balance invariant
3. A way to implement the insert, remove and transfer operations.

Sized Lists

The first part is super easy. Let's define a type:

```
data SList a = SL { size :: Int, elems :: [a] }
```

We have a special field that saves the size because otherwise, we have a linear time computation that wrecks Okasaki's careful analysis. (Actually, he presents a variant which does *not* require saving the size as well, but that's for another day.)

Think Aloud:

Read the question aloud and voice your thoughts while solving the exercise.

How can we be sure that `size` is indeed the *real size* of `elems`?
Write a measure `realSize` to get the number of elements in any list:

```
-- write measure realSize in here
```

Hint: When you are done, uncomment `data SList`, that specifies a *refined* type for `SList` that ensures that the *real* size is saved in the `size` field.

```
-- {-@ data SList a = SL {
--     size  :: Nat
--     , elems :: {v:[a] | realSize v = size}
--     }
-- @-}
```

Answer

```
{-@ measure realSize @-} realSize      :: [a] -> Int
realSize []      = 0
realSize (_,xs) = 1 + realSize xs
```

Now, as a sanity check, consider this:

```
okList  = SL 1 ["cat"]    -- accepted
badList = SL 1 []         -- rejected
```

LET'S DEFINE AN ALIAS for lists of a given size `N`:

```
{-@ type SListN a N = {v:SList a | size v = N} @-}
```

Think Aloud: Read the question aloud and voice your thoughts while solving the exercise.

NOW DEFINE AN ALIAS `NEList` for lists that are not empty by replacing `{true}` by the correct signature.

```
{-@ type NEList a = {true} @-}
```

Answer

```
{-@ type NEList a = {v:SList a | size v > 0} @-}
```

Finally, we can define a basic API for `SList`.

To CONSTRUCT LISTS, we use `nil` and `cons`:

```
{-@ nil :: SListN a 0 @-}
nil = SL 0 []

{-@ cons :: a -> xs:SList a -> SListN a {size xs + 1} @-}
cons x (SL n xs) = SL (n+1) (x:xs)
```

Think Aloud: Read the question aloud and voice your thoughts while solving the exercise.

Exercise 3.1 (Destructing Lists). *We can destruct lists by writing a `hd` (head) and `tl` (tail) function as shown below.*

For `tl`, fix the signature such that it receives a non-empty list and returns another without the first element.

For `hd`, do the opposite. From the presented signature, write the implementation. This function returns just the element at the front of the list.

```
{-@ tl      :: {xs:NEList a | size xs < 0}
              -> SListN a {size xs} @-}
tl (SL n (_:xs)) = SL (n-1) xs
tl _             = die "empty SList"

{-@ hd'     :: xs:NEList a -> a @-}
hd' _ = die "empty SList"
```

Hint: When you are done, `okHd` should be verified, but `badHd` should be rejected.

```
{-@ okList :: SListN String 1 @-}

okHd = hd' okList      -- accepted

badHd = hd' (tl okList) -- rejected
```

Answer

```
{-@ tl      :: xs:NEList a -> SListN a {size xs - 1}
@-} hd' (SL _ (x:_)) = x hd' _ = die "empty SList"
```

Queue Type

It is quite straightforward to define the Queue type, as a pair of lists, front and back, such that the latter is always smaller than the former:

```

{-@ data Queue a = Q {
    front :: SList a
    , back  :: SListLE a (size front)
  }
  @-}
data Queue a = Q
  { front :: SList a
  , back  :: SList a
  }

```

THE ALIAS `SListLE a L` corresponds to lists with at most `N` elements:

```

{-@ type SListLE a N = {v:SList a | size v <= N} @-}

```

As a quick check, notice that we *cannot represent illegal Queues*:

```

okQ  = Q okList nil  -- accepted, |front| > |back|
badQ = Q nil okList  -- rejected, |front| < |back|

```

Queue Operations

Almost there! Now all that remains is to define the Queue API. The code below is more or less identical to Okasaki's (I prefer `front` and `back` to his `left` and `right`.)

THE EMPTY QUEUE is simply one where both `front` and `back` are both empty:

```

emp = Q nil nil

```

Think Aloud: Read the question aloud and voice your thoughts while solving the exercise.

Exercise 3.2 (Queue Sizes). *For the remaining operations we need some more information. Do the following steps:*

1. Write a measure `qsize` to describe the queue size,
2. Use it to complete the definition of `QueueN` below replacing the `true` predicate, and
3. Check the use of the alias in the examples.

```
-- | create measure qsize here

-- | Queues of size `N`
{-@ type QueueN a N = {v:Queue a | true} @-}

{-@ emp :: QueueN _ 0 @-}

{-@ example2Q :: QueueN _ 2 @-}
example2Q = Q (1 `cons` (2 `cons` nil)) nil

{-@ example0Q :: QueueN _ 0 @-}
example0Q = Q nil nil
```

Answer

```
{-@ measure qsize @-} qsize      :: Queue a -> Int qsize
(Q l r) = size l + size r {-@ type QueueN a N = {v:Queue a |
qsize v = N} @-}
```

To REMOVE an element we pop it off the front by using `hd` and `tl`. Notice that the `remove` is only called on non-empty Queues, which together with the key balance invariant (makeq that we will see later), ensures that the calls to `hd` and `tl` are safe.

Think Aloud: Read the question aloud and voice your thoughts while solving the exercise.

Add a LiquidHaskell signature to `remove` using `QueueN`. When you are done, `okRemove` should be accepted, `badRemove` should be rejected.

```
remove (Q f b) = (hd f, makeq (tl f) b)

okRemove = remove example2Q  -- accept
badRemove = remove example0Q -- reject
```

Answer

```
{-@ remove :: {q:Queue a | qsize q > 0} -> (a, QueueN a
{qsize q - 1}) @-}
```

To INSERT an element we just `cons` it to the back list, and call the *smart constructor* `makeq` to ensure that the balance invariant holds.

Think Aloud: Read the question aloud and voice your thoughts while solving the exercise.

Exercise 3.3 (Insert). *Write down a liquid type signature for replicate (that uses insert), so that it adds the same element n times to the queue, and okReplicate is accepted by LiquidHaskell, but badReplicate is rejected.*

```
{-@ insert :: a -> q:Queue a -> QueueN a {qsize q + 1}  @-}
insert e (Q f b) = makeq f (e `cons` b)

-- write liquid type signature
replicate 0 _ = emp
replicate n x = insert x (replicate (n-1) x)

{-@ okReplicate :: QueueN _ 3 @-}
okReplicate = replicate 3 "Yeah!" -- accept

{-@ badReplicate :: QueueN _ 3 @-}
badReplicate = replicate 1 "No!" -- reject
```

Answer

```
{-@ replicate :: n:Nat -> a -> QueueN a n @-}
```

TO ENSURE THE INVARIANT we use the smart constructor makeq, which is where the heavy lifting happens. The constructor takes two lists, the front f and back b and if they are balanced, directly returns the Queue, and otherwise transfers the elements from b over using the rotate function rot described next. However, rot still needs to have an appropriate liquid signature.

```
{-@ makeq :: f:SList a ->
      b:SListLE a {size f + 1 } ->
      QueueN a {size f + size b} @-}
makeq f b
  | size b <= size f = Q f b
  | otherwise       = Q (rot f b nil) nil
```

THE ROTATE FUNCTION will ensure that the two lists are balanced, as introduced at the start. It is arranged so that it the hd is built up fast, before the entire computation finishes; which, combined with laziness provides the efficient worst-case guarantee.

Think Aloud: Read the question aloud and voice your thoughts while solving the exercise.




Figure 3.4: Transferring Elements from back to front.

Exercise 3.4 (Rotate). *Read and fix the liquid type added to `rot`, following the next properties.*

The Rotate function `rot`:

1. *is only called when back is one larger than the front (we never let things drift beyond that).*
2. *And the return size is the sum of the size in front, back and the additional to be rotated.*

```
{-@ rot :: f:SList a
    -> b:SListN a {1 - size f}
    -> acc:SList a
    -> SListN a {size acc}
@-}
rot f b acc
  | size f == 0 = hd b `cons` acc
  | otherwise = hd f `cons` rot (tl f) (tl b) (hd b `cons` acc)
```

Answer

```
{-@ rot :: f:SList a -> b:SListN a {1 + size f} -> acc:SList
a -> SListN a {size f + size b + size acc} @-}
```

Recap of everything!

Well there you have it; Okasaki's beautiful lazy Queue, with the invariants easily expressed and checked with LiquidHaskell. This example is particularly interesting because

1. The refinements express invariants that are critical for efficiency,
2. The code introspects on the size to guarantee the invariants, and
3. The code is quite simple and we hope, easy to follow!

This exercise concludes the Short Tutorial of LiquidHaskell. Thank you for tagging along!

4

Cheat Sheet

Welcome to the LiquidHaskell Short Tutorial Cheat Sheet!

Here are the main concepts and examples you can use to complete the exercises.

Specifications

LiquidHaskell specifications in functions are written between `{-@ spec @-}`.

For example:

```
{-@ calcPer    :: {a:Int | a > 0} -> {b:Int | 0 <= b && b <= a} -> c:Int @-}
calcPer      :: Int -> Int -> Int
calcPer a b   = (b * 100) `div` a
```

Alias

To reuse a specification we can use aliases.

For example:

```
{-@ type Nat    = {v:Int | 0 <= v}      @-}
{-@ type Btwn Lo Hi = {v:Int | Lo <= v && v < Hi} @-}
```

Liquid types in Datatypes

To add a specification to the datatypes, first create the datatype in Haskell and then add the specification inside `{-@ @-}`.

For example:

1) In Haskell

```
data Sparse a = SP { spDim    :: Int
                    , spElems :: [(Int, a)] }
```

2) Adding the LiquidHaskell specification

```
{-@ data Sparse a = SP { spDim    :: Nat
                        , spElems :: [(Btwn 0 spDim, a)] } @-}
```

Measures

Measures lift an Haskell function to the refinements logic. It is first created as a Haskell function, sinalizing that it is a measure and adding liquid types to the signature. Then, it can be used inside other refinements.

For example:

```
{-@ measure mySize @-}
{-@ mySize :: [a] -> Nat @-}
mySize :: [a] -> Int
mySize []      = 0
mySize (_:rs) = 1 + mySize rs
```

And then, size can be used in:

```
{-@ type ListN a N = {v:[a] | size v == N} @-}
```


5

Refined Datatypes

So far, we have seen how to refine the types of *functions*, to specify, for example, pre-conditions on the inputs, or post-conditions on the outputs. In this section we will see how to apply these in datatypes. First, by defining properties of data values, and then by defining *datatypes* that satisfy certain invariants. In the latter, it is handy to be able to directly refine the data definition, making it impossible to create illegal inhabitants.

MEASURES are used to define *properties* of Haskell data values that are useful for specification and verification.

A MEASURE is a *total* Haskell function, 1. With a *single* equation per data constructor, and 2. Guaranteed to *terminate*, typically via structural recursion.

We can tell LiquidHaskell to *lift* a function meeting the above requirements into the refinement logic by declaring:

```
{-@ measure nameOfMeasure @-}
```

For example, for a list we can define a way to *measure* its size with the following function.

```
{-@ measure size @-}  
{-@ size :: [a] -> Nat @-}  
size :: [a] -> Int  
size []      = 0  
size (_:rs) = 1 + size rs
```

Then, we can use this measure to define aliases.

Let's create another measure named `notEmpty` that takes a list as input and returns a `Bool` with the information if it is empty or not.

```
-- write notEmpty measure
```

Answer

```
{-@ measure notEmpty @-} notEmpty      :: [a] -> Bool
notEmpty []      = False notEmpty (_,_) = True
```

We can now define a couple of useful aliases for describing lists of a given dimension.

For example, we can define that a list has exactly N elements.

```
{-@ type ListN a N = {v:[a] | size v == N} @-}
```

Note that when defining refinement type aliases, we use uppercase variables like N to distinguish *value* parameters from the lowercase *type* parameters like a .

Now, try to create an alias `NEList` for an empty list, using the measure `notEmpty` created before. When removed from comment, the first example should raise an error while the second should not.

```
-- write the alias here
```

```
-- {-@ ne1 :: NEList Int@-}
-- ne1 = [] :: [Int]
-- {-@ ne1 :: NEList Int@-}
-- ne2 = [1,2,3,4] :: [Int]
```

Answer

```
{-@ type NEList a = {v:[a] | notEmpty v} @-}
```

Sparse Vectors

As our first example of a refined datatype, let's see Sparse Vectors. While the standard `Vector` is great for dense arrays, often we have to manipulate sparse vectors where most elements are just 0. We might represent such vectors as a list of index-value tuples `[(Int, a)]`.

Let's create a new datatype to represent such vectors:

```
data Sparse a = SP { spDim    :: Int
                    , spElems :: [(Int, a)] }
```

Thus, a sparse vector is a pair of a dimension and a list of index-value tuples. Implicitly, all indices *other* than those in the list have the value 0 or the equivalent value type a .

LEGAL

Sparse vectors satisfy two crucial properties. 1. the dimension stored in `spDim` is non-negative;

2. every index in `spElems` must be valid, i.e. between 0 and the dimension.

Unfortunately, Haskell's type system does not make it easy to ensure that *illegal vectors are not representable*.

DATA INVARIANTS LiquidHaskell lets us enforce these invariants with a refined data definition:

```
{-@ data Sparse a = SP { spDim    :: Nat
                        , spElems :: [(Btwn 0 spDim, a)] } @-}
```

Where, as before, we use the aliases:

```
{-@ type Nat      = {v:Int | 0 <= v}      @-}
{-@ type Btwn Lo Hi = {v:Int | Lo <= v && v < Hi} @-}
```

REFINED DATA CONSTRUCTORS The refined data definition is internally converted into refined types for the data constructor `SP`. So, by using refined input types for `SP` we have automatically converted it into a *smart* constructor that ensures that *every* instance of a `Sparse` is legal. Consequently, LiquidHaskell verifies:

```
okSP :: Sparse String
okSP = SP 5 [ (0, "cat")
              , (3, "dog") ]
```

but rejects, due to the invalid index:

```
badSP :: Sparse String
badSP = SP 5 [ (0, "cat")
               , (6, "dog") ]
```

Write another example of a `Sparse` data type that is invalid. Remove the comment from the type signature below and complete the implementation with the example.

```
-- badSP' :: Sparse String
```

Answer

e.g., `badSP' = SP -1 [(0, "cat")]`

FIELD MEASURES It is convenient to write an alias for sparse vectors of a given size N . So that we can easily say in a refinement that we have a sparse vector of a certain size.

For this we can use *measures*.

MEASURES WITH SPARSE VECTORS

Similarly, the sparse vector also has a *measure* for its dimension, but in this case it is already defined by `spDim`, so we can use it to create the new alias of sparse vectors of size N .

Think Aloud:

For the following exercise, we will use a technique called Think Aloud, where you should try to say everything that comes to your mind while you engage with the exercise.

In specific, aim:

- a) to speak all thoughts, even if they are unrelated to the task;
- b) to refrain from explaining the thoughts;
- c) to not try to plan out what to say;
- d) to imagine that you are alone and speaking to yourself; and
- e) to speak continuously.

For the following exercise, read the question aloud and remember to voice your thoughts while solving the exercise.

Following what we did with the lists, write the alias `SparseN` for sparse vector of length N , using `spDim` instead of `size`.

Hint: When you are done, you can see how we can use `SparseN` in the example below.

```
-- write the alias here
```

Answer

e.g., `{-@ type SparseN a N = {v: Sparse a | spDim v == N} @-}`

Sparse Products

Vectors are similar to Sparse Vectors, and therefore, have a *measure* of size named `vlen`. So, now, we can see that LiquidHaskell is able to compute a sparse product, making the product of all the same indexes and returning its sum. Remove the comments and run the code ahead.

```
-- dotProd :: Vector Int -> Sparse Int -> Int
-- {-@ dotProd :: x:Vector Int -> SparseN Int (vlen x) -> Int @-}
-- dotProd x (SP _ y) = go 0 y
-- where
--   go sum ((i, v) : y') = go (sum + (x ! i) * v) y'
--   go sum []             = sum
```

LiquidHaskell verifies the above by using the specification to conclude that for each tuple (i, v) in the list `y`, the value of `i` is within the bounds of the vector `x`, thereby proving `x ! i` safe.

YOU FINISHED THE SECOND PART OF THE TUTORIAL!

You finished the Tutorial! Tell the interviewers you got to the end of the page, and answer some questions from our team before moving to the next section.

NEXT EXERCISE!

Now that you have learned the main blocks of LiquidHaskell, let's complete an exercise using all the concepts.

You can open a Cheat Sheet with examples of the main concepts on another tab on the side.

Next